

Industrial Project

2026-07-11

No cover image available.

Engineering Work Order Profit and Loss Analytics: Synthetic Data Platform for Revenue, Cost, Margin, Billing, and Risk Monitoring

A synthetic-data-based analytics platform for estimating work-order revenue, tracking delivery-center and field-operations labor, monitoring billing and other costs, calculating profit and loss, and identifying margin risk across the engineering work-order lifecycle.

Project Snapshot

DATE

2026-07-11

CATEGORY

Industrial Project

SHORT DESCRIPTION

A synthetic-data-based analytics platform for estimating work-order revenue, tracking delivery-center and field-operations labor, monitoring billing and other costs, calculating profit and loss, and identifying margin risk across the engineering work-order lifecycle.

TAGS / TECH STACK

Data Analytics

Financial Analytics

Profit and Loss

Work Order Management

Streamlit

Snowflake

Python

Plotly

Synthetic Data

Risk Monitoring

Decision Support

Data Quality

Project Links

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

Project Link

GitHub Link

GitHub Link

Project Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

GitHub Link

Full Project Content

Engineering Work Order Profit and Loss Analytics is a Streamlit and Snowflake-ready analytics prototype for estimating work-order revenue, tracking labor and operational costs, monitoring billing, calculating profit and loss, and identifying margin risk across the engineering delivery lifecycle.

[View the GitHub Repository.](#)

Wiki documentation: [Home](#) · [Getting Started](#) · [Architecture](#) · [Data Model](#) · [Dashboard Guide](#)

Application documentation: [Synthetic Data](#) · [Financial Calculations](#) · [Snowflake Integration](#) · [Deployment](#)

Engineering and governance: [Testing and Quality](#) · [Security and Privacy](#) · [Troubleshooting](#) · [Roadmap](#)

One-Line Idea

Turn fragmented engineering work-order revenue, labor, billing, and cost records into a single explainable profit-and-loss monitoring layer—without using confidential production data in the public repository.

Why This Project Exists

Engineering delivery teams often manage work orders through multiple operational and financial signals. Revenue may be represented through work units or sales-order lines. Labor effort may be recorded through activity or daily reports. Billing may arrive later through invoice or progress-billing events. Travel, permit, subcontractor, and exception costs may sit in separate sources.

This creates a visibility gap. A team may know that a work order is active, but not immediately know:

- how much revenue is currently forecast;
- how much delivery-center and field-operations labor has accumulated;
- whether billing is lagging behind operational progress;
- whether travel, rework, permitting, or other costs are consuming margin;
- which work orders are below target margin;
- which projects should be reviewed before they become losses.

This project organizes those signals into a synthetic, reproducible analytics platform that can be demonstrated publicly and later mapped to governed enterprise data sources.

The goal is not to create an audited accounting system. The goal is to create a practical operational analytics layer that helps engineering, operations, finance, and delivery leaders understand work-order profitability earlier in the lifecycle.

Project at a Glance

Area	Description
Project type	Industrial analytics and decision-support prototype
Domain	Engineering work-order delivery, cost monitoring, billing visibility, and margin risk
Data mode	Fully synthetic public dataset
Primary inputs	Work orders, work-unit revenue, activity/labor events, billing events, other costs, labor rates, milestone mapping
Main output	One-row-per-work-order financial and lifecycle summary
Dashboard	Streamlit application with Plotly charts and downloadable summaries
Data warehouse path	Snowflake-ready schemas, analytical views, and quality checks
Core language	Python
Main libraries	pandas, Streamlit, Plotly, pytest
Documentation	README, architecture guide, data dictionary, root-level technical guides, and GitHub Wiki pages
Current maturity	Public synthetic-data MVP and reference architecture
Safety boundary	Demonstration only; not an audited finance, revenue-recognition, or accounting system

What Makes the Platform Different

The project treats work-order profit and loss as both a **financial analytics problem** and a **delivery lifecycle problem**.

Conventional approach	Platform approach
Review revenue, billing, and labor separately	Consolidate all signals into a work-order summary
Wait until late-stage billing to understand margin	Estimate margin throughout the lifecycle
Track labor totals without lifecycle context	Break effort down by milestone, activity type, and delivery organization

Conventional approach	Platform approach
Look at financial performance only after completion	Flag high-risk work orders while they are still actionable
Treat rate assumptions as fixed code	Store rates as configurable data for scenario analysis
Ignore zero-hour lifecycle events	Retain them as operational milestones while excluding them from labor cost
Use confidential extracts for demos	Use deterministic synthetic data for public reproducibility
Build a dashboard without data-quality checks	Include duplicate, orphan, missing, zero-hour, and negative-hour checks


Important: The repository is intentionally generic. It does not contain confidential customer, employee, operational, financial, or production data. All records, identifiers, names, rates, dates, and values are synthetic.

Reader Walkthrough

A reader can understand the platform in nine steps:

1. Synthetic work-order master records define account, program, region, job category, status, intake date, and margin target.
2. Work-unit revenue lines provide estimated revenue using quantity and unit-rate logic.
3. Activity/labor records capture milestone, activity type, delivery organization, date, and hours.
4. Billing events represent invoice or progress-billing activity.
5. Other-cost records capture travel, permit, subcontractor, or miscellaneous expenses.
6. Labor-rate assumptions convert hours into delivery-center and field-operations labor cost.
7. The metric engine produces revenue, cost, margin, billed amount, variance, travel, rework, and risk reasons.
8. The Streamlit dashboard exposes portfolio KPIs, drill-downs, lifecycle charts, rate scenarios, and quality checks.
9. Snowflake SQL scripts show how the same logical model can be implemented over governed warehouse tables or views.

Platform Architecture



The platform is organized around two complementary implementation paths.

1. Synthetic Analytics Application

The public implementation loads deterministic synthetic CSV files, validates and standardizes the records, combines work-unit revenue, DWR labor, billing, and other-cost data, and produces a one-row-per-work-order financial summary. The Streamlit application presents portfolio KPIs, work-order drill-downs, lifecycle and labor analysis, rate scenarios, data-quality findings, risk reasons, and downloadable summaries.

2. Governed Snowflake Deployment

The production path replaces local CSV files with approved Snowflake tables or secure views. Snowflake-compatible schemas and analytical views support governed source mapping, configurable labor rates and thresholds, validated financial calculations, role-based access, and deployment through Streamlit in Snowflake or another approved analytics layer.

[Read the Architecture Wiki page](#)

Data Model

The project uses a compact star-style analytical model.

Table	Grain	Purpose
dim_work_order	One row per work order	Master work-order attributes and target margin

Table	Grain	Purpose
fact_work_unit_revenue	One row per work-unit revenue line	Estimated revenue from work-unit quantity and rate
fact_dwr_labor	One row per DWR labor or lifecycle event	Hours by milestone, activity type, and delivery organization
fact_billing	One row per billing event	Invoice, progress billing, or billing status events
fact_other_cost	One row per non-labor cost event	Travel, permit, subcontractor, or miscellaneous costs
dim_labor_rate	One row per delivery organization/rate	Burdened hourly cost assumptions
dim_milestone	One row per milestone/activity mapping	Lifecycle ordering and standardized activity grouping

The central analytical output is a one-row-per-work-order summary used by the dashboard.

[Read the Data Model Wiki page](#)

Synthetic Dataset

The bundled dataset is deterministic and can be regenerated with the default seed.

Dataset	Rows	Analytical use
Work orders	600	Portfolio, filters, status, target margin, account/program/region/category
Work-unit revenue	2,660	Estimated revenue and sales-order/work-unit drill-down
DWR activity/labor records	8,713	Labor hours, lifecycle effort, delivery organization, travel, rework
Billing events	730	Billed amount, billing progress, billing visibility
Other costs	977	Travel, permit, subcontractor, and operational cost tracking
Labor rates	2	Delivery-center and field-operations rate assumptions
Milestone mapping	16	Lifecycle ordering and dashboard grouping

The dataset spans January 2025 through May 2026 and is designed only for demonstration. It should not be interpreted as an industry benchmark.

[Read the Synthetic Data Wiki page](#)

Financial Calculation Engine

The core calculations are intentionally transparent.

```
Estimated Revenue = Σ(Work Unit Quantity × Item Rate)
Labor Cost = Σ(Activity Hours × Applicable Hourly Rate)
Actual Cost = Labor Cost + Other Cost
Forecast Revenue = MAX(Estimated Revenue, Billed Revenue)
Forecast Margin = Forecast Revenue - Actual Cost
Forecast Margin % = Forecast Margin ÷ Forecast Revenue
Margin Variance % = Forecast Margin % - Target Margin %
```

Risk bands are rule-based and interpretable.

Risk band	Example trigger
High	Margin below high-risk threshold, high rework, or high travel
Medium	Margin below medium-risk threshold, moderate rework, or moderate travel
Low	No configured risk condition is met

The platform exposes risk reasons instead of hiding them inside an opaque score. This makes the dashboard easier to review with operations and finance stakeholders.

[Read the Financial Calculations Wiki page](#)

Dashboard Experience

The Streamlit dashboard is organized around practical questions.

Portfolio Overview

Shows total revenue, cost, profit, margin, billed amount, risk distribution, program-level performance, regional performance, and a downloadable work-order summary.

Work Order Detail

Provides a drill-down into one selected work order, including work-unit revenue, billing history, other costs, lifecycle activity, labor hours, margin, and risk reason.

Lifecycle and Labor

Compares hours across milestones and delivery organizations. This helps identify where effort is being consumed across intake, planning, field operations, design, permitting, delivery, and billing.

Rate Scenario

Allows users to change delivery-center and field-operations hourly rates and immediately observe how profit, margin, and risk counts change.

Data Quality

Surfaces potential duplicate work-unit rows, orphan activity records, zero-hour lifecycle events, negative hours, and work orders without revenue lines.

[Read the Dashboard Guide](#)

Snowflake-Ready Integration

The implementation includes Snowflake-oriented SQL inside the application folder.

Repository file	Purpose
<code>insite_pl_mvp/sql/snowflake_schema.sql</code>	Defines Snowflake tables for the logical work-order model
<code>insite_pl_mvp/sql/mvp_views.sql</code>	Builds analytical views for revenue, labor, billing, cost, margin, and work-order reporting
<code>insite_pl_mvp/tests/test_metrics.py</code>	Validates core metric behavior; application-level data-quality findings are displayed in Streamlit

A production implementation should not upload confidential data into a public repository. Instead, it should map approved enterprise source fields into curated warehouse tables or secure views and connect the dashboard to those governed objects.

[Read the Snowflake Integration Wiki page](#)

Technology Stack

Layer	Technologies
Application	Streamlit
Language	Python

Layer	Technologies
Data processing	pandas
Visualization	Plotly
Testing	pytest and local validation
Data warehouse path	Snowflake SQL
Documentation	Markdown, README, docs, GitHub Wiki
Data	Synthetic CSV files generated by Python
Packaging	requirements.txt and setup scripts

Repository Structure

```
Engineering-Work-Order-Profit-and-Loss-Analytics/  
├── assets/  
│   └── diagrams/  
│       └── workorder_profit_loss_architecture.svg  
├── insite_pl_mvp/  
│   ├── data/  
│   │   ├── dim_labor_rate.csv  
│   │   ├── dim_milestone.csv  
│   │   ├── dim_work_order.csv  
│   │   ├── fact_billing.csv  
│   │   ├── fact_dwr_labor.csv  
│   │   ├── fact_other_cost.csv  
│   │   └── fact_work_unit_revenue.csv  
│   ├── scripts/  
│   │   └── generate_synthetic_data.py  
│   ├── sql/  
│   │   ├── mvp_views.sql  
│   │   └── snowflake_schema.sql  
│   ├── src/  
│   │   ├── __init__.py  
│   │   ├── data_loader.py  
│   │   └── metrics.py  
│   ├── tests/  
│   │   └── test_metrics.py  
│   ├── DATA_DICTIONARY.md  
│   ├── README.md  
│   ├── app.py  
│   ├── requirements.txt  
│   └── setup_and_run.bat
```

```
| └─ setup_and_run.ps1
|─ A-Z_FILE_GUIDE.md
|─ ARCHITECTURE.md
|─ CHANGELOG.md
|─ CODE_OF_CONDUCT.md
|─ CONTRIBUTING.md
|─ DATA_DICTIONARY.md
|─ GITHUB_REPOSITORY_SETUP.md
|─ PROJECT_STRUCTURE.md
|─ README.md
|─ ROADMAP.md
|─ SECURITY.md
|─ configuration.md
|─ data-model.md
|─ deployment.md
|─ developer-guide.md
|─ faq.md
|─ financial-calculations.md
|─ installation.md
|─ security-privacy.md
|─ snowflake-integration.md
|─ synthetic-data.md
|─ testing.md
|─ troubleshooting.md
|─ user-guide.md
```

[Read the Project Structure guide](#)

Local Run

Clone the repository and enter the application folder:

```
git clone https://github.com/dranubhaparashar/Engineering-Work-Order-Profit-and-Loss-
Analytics.git
cd Engineering-Work-Order-Profit-and-Loss-Analytics\insite_pl_mvp
```

Windows one-click setup

```
.\setup_and_run.ps1
```

Alternatively, double-click:

```
setup_and_run.bat
```

Windows manual setup

```
py -m venv .venv
.\.venv\Scripts\Activate.ps1
python -m pip install --upgrade pip
pip install -r requirements.txt
streamlit run app.py
```

macOS or Linux

```
git clone https://github.com/dranubhaparashar/Engineering-Work-Order-Profit-and-Loss-
Analytics.git
cd Engineering-Work-Order-Profit-and-Loss-Analytics/insite_pl_mvp
python3 -m venv .venv
source .venv/bin/activate
python -m pip install --upgrade pip
pip install -r requirements.txt
streamlit run app.py
```

Open the Streamlit URL shown in the terminal, normally:

```
http://localhost:8501
```

[Read Installation](#)

Current Implementation Status

Capability	Status
Synthetic dataset generation	Implemented
Streamlit dashboard	Implemented
Work-order financial summary	Implemented
Work-unit revenue aggregation	Implemented
Labor-cost calculation	Implemented
Other-cost integration	Implemented
Billing visibility	Implemented
Margin and risk classification	Implemented

Capability	Status
Rate scenario analysis	Implemented
Data-quality page	Implemented
Snowflake schema and views	Implemented
Automated pytest coverage	Implemented (insite_pl_mvp/tests/test_metrics.py)
GitHub Actions workflow	Not currently present in the live repository
GitHub Wiki content	Included
Enterprise authentication	Not included in public demo
Audited finance/accounting controls	Not included
Real production data integration	Intentionally not included

Known Limitations

The project is designed as a public, synthetic demonstration. Important limitations include:

1. The financial logic is not an accounting standard and should not be treated as audited finance logic.
2. Revenue recognition, overhead allocation, tax, cost capitalization, and labor-burden rules vary by organization.
3. The risk model is rules-based rather than predictive or statistically validated.
4. CSV loading is appropriate for an MVP; production usage should rely on governed warehouse views.
5. The bundled dataset is synthetic and should not be interpreted as an industry performance benchmark.
6. Authentication, authorization, row-level security, audit logging, and secrets management must be added for real deployments.
7. Any mapping to enterprise data must be approved by the data owner and reviewed for privacy, pricing, contract, employee, and customer sensitivity.

These limitations are explicitly documented so the portfolio post does not overclaim the maturity or authority of the platform.

Roadmap

Phase 1 — Public Synthetic MVP

- Provide a complete synthetic dataset.
- Build a Streamlit dashboard.
- Add profit-and-loss calculations.
- Add rate scenarios and data-quality checks.
- Document the repository and Wiki.

Phase 2 — Warehouse-Backed Analytics

- Replace CSV reads with Snowflake views.
- Add configuration for environment-specific sources.
- Materialize the work-order summary model.
- Add row-level access controls and secure roles.
- Validate field mappings with finance and operations stakeholders.

Phase 3 — Operational Governance

- Add authentication and authorization.
- Add audit logging for exports and scenario assumptions.
- Define approved margin thresholds and exception workflows.
- Add reviewer notes, dispositions, and follow-up ownership.

Phase 4 — Advanced Intelligence

- Add anomaly detection for unusual labor, travel, rework, or billing lag.
- Add forecasting for likely final margin.
- Add explainable drivers for margin deterioration.
- Add project-completion and cash-flow forecasting.

[Read the Roadmap](#)

Security, Privacy, and Claim Boundary

Work-order financial analytics can involve sensitive information, including customer names, employee identifiers, labor rates, contract terms, billing values, site locations, operational notes, and internal margin targets.

The public repository therefore follows a strict boundary:

- no confidential company data;
- no customer or employee records;
- no production work-order numbers;
- no real billing or pricing extracts;
- no proprietary internal source files;
- no credentials or secrets;

- no claim of audited financial correctness.

Before connecting the project to real systems, teams should obtain written approval, use role-based access control, mask sensitive fields, keep secrets outside source control, validate all formulas, and review repository history before publication.

[Read Security and Privacy.](#)

Testing and Validation Strategy

The repository includes automated tests for data quality and financial logic. A stronger production test plan should add:

- source-to-target reconciliation tests;
- revenue aggregation tests;
- labor-rate effective-date tests;
- billing status tests;
- margin-threshold boundary tests;
- duplicate and orphan record tests;
- large-data performance tests;
- dashboard regression tests;
- Snowflake view validation;
- security and access-control tests;
- finance stakeholder sign-off.

[Read Testing and Quality.](#)

Complete Wiki Documentation

Wiki page	Purpose
Wiki Home	Project entry point and documentation navigation
Getting Started	Local setup and first run
Architecture	System layers, data flow, and deployment patterns
Data Model	Dimensions, facts, grain, and join logic
Dashboard Guide	Dashboard pages and user workflow
Synthetic Data	Generator, fields, assumptions, and safe public data policy
Financial Calculations	Revenue, cost, margin, billing, and risk formulas

Wiki page	Purpose
Snowflake Integration	Warehouse schema, views, and production mapping
Deployment	Local, Streamlit, and Snowflake deployment options
Testing and Quality	Automated checks and production validation plan
Troubleshooting	Common setup and runtime issues
Security and Privacy	Safe public-use boundary and real-data controls
FAQ	Common user and developer questions
Roadmap	Future development phases

Practical Use Cases

User group	How the platform helps
Engineering managers	Monitor work-order margin risk across active delivery work
Operations leaders	Identify where labor, travel, or rework is increasing cost
Finance analysts	Compare forecast revenue, billed revenue, cost, and margin
Delivery-center teams	Understand lifecycle effort and milestone-level workload
Field-operations teams	Review travel-heavy or field-hour-heavy work orders
Data engineers	Use the logical model to design Snowflake views
BI developers	Adapt the metric layer into dashboards or reporting systems
Governance teams	Review synthetic-data, security, privacy, and claim boundaries
Portfolio reviewers	Evaluate project performance without exposing confidential data

Key Strengths

- Uses fully synthetic public data.
- Demonstrates a complete work-order analytics workflow.
- Combines revenue, labor, billing, other costs, and risk into one model.
- Preserves drill-down detail while also producing portfolio KPIs.
- Provides clear formulas instead of opaque financial logic.

- Separates delivery-center and field-operations labor assumptions.
 - Includes travel and rework risk indicators.
 - Includes data-quality checks inside the application.
 - Provides Snowflake-ready schema and analytical views.
 - Includes README, architecture documentation, data dictionary, implementation guides, and GitHub Wiki pages.
 - Avoids company-specific, customer-specific, or source-system-specific public claims.
-

Key Innovation

The strongest contribution is the translation of scattered work-order operational and financial events into a transparent, explainable, and privacy-safe profit-and-loss monitoring layer.

The project shows how a team can move from late-stage financial visibility to lifecycle-based margin awareness. Instead of waiting until final billing, stakeholders can review estimated revenue, accumulated labor, non-labor costs, billing progress, target-margin variance, and risk reasons while the work order is still active.

Project Links

- [GitHub Repository](#)
 - [GitHub Wiki](#)
 - [README](#)
 - [Application Folder](#)
 - [Architecture](#)
 - [Data Dictionary](#)
 - [Financial Calculations](#)
 - [Snowflake Integration](#)
 - [Streamlit App](#)
 - [Metric Engine](#)
 - [Synthetic Data Generator](#)
 - [Snowflake Schema](#)
 - [Analytical Views](#)
 - [Tests](#)
 - [Testing and Quality](#)
 - [Security and Privacy](#)
 - [Roadmap](#)
-

Conclusion

Engineering Work Order Profit and Loss Analytics demonstrates how synthetic data, transparent calculations, and a lightweight dashboard can turn fragmented delivery and financial events into actionable work-order intelligence.

The platform estimates revenue from work-unit lines, prices labor through configurable rate assumptions, adds operational costs, monitors billing events, calculates forecast margin, and surfaces risk reasons that stakeholders can understand. Its value is strongest as a reproducible public MVP and reference architecture for teams that want to build safer, governed work-order financial analytics without exposing confidential data.

The next step for a real organization would be controlled source mapping, finance-approved formulas, warehouse-level governance, access control, audit logging, and validation against approved historical records.

Final Thought

Estimate early. Track continuously. Explain margin risk. Protect sensitive data.

Work-order revenue · labor cost · billing visibility · profit and loss · margin risk · synthetic data · Streamlit · Snowflake-ready analytics · complete GitHub Wiki