


Industrial Project

2026-04-23

# Execution-Aware Agentic VRP Solver and Benchmark Studio



Production-grade Vehicle Routing Problem solver with 8 solver backends, 18 operational scenarios, and an Agentic AI layer powered by Google Gemini 2.0 Flash.

## Project Snapshot

### DATE

2026-04-23

### CATEGORY

Industrial Project

### SHORT DESCRIPTION

Production-grade Vehicle Routing Problem solver with 8 solver backends, 18 operational scenarios, and an Agentic AI layer powered by Google Gemini 2.0 Flash.

### TAGS / TECH STACK

AI

VRP

Vehicle Routing

OR-Tools

PyVRP

Agentic AI

Gemini

Streamlit

Optimization

MLOps

Hugging Face

## Project Links

[Demo Video](#)

[Hugging Face Space](#)

[GitHub Link](#)

[Project Link](#)

[GitHub Link](#)

[GitHub Link](#)

[Demo Video](#)

## Full Project Content

Building an AI system that doesn't just **solve routes** — but combines 8 solver backends, 18 real-world scenarios, an Agentic AI orchestrator, and a benchmark-grade Streamlit studio into a single execution-aware optimization platform.

**GitHub Repository:** [dranubhaparashar/Execution-Aware-Agentic-VRP-Solver-and-Benchmark-Studio](https://github.com/dranubhaparashar/Execution-Aware-Agentic-VRP-Solver-and-Benchmark-Studio)

 **Live demo video:** [youtu.be/FqVuVjW20yo](https://youtu.be/FqVuVjW20yo)

 **Try it live:** [Hugging Face Space](#)

## Vision

Most VRP solvers stop at producing a route plan.

This project focuses on the full operational chain required to make routing decisions usable in a **live dispatch environment** where technicians are already in motion, jobs get cancelled mid-day, urgent orders inject themselves into the schedule, and shifts change without notice:

- Ingest 40,000 orders, 400 routes, 8 depots, multi-customer goal profiles
- Run 8 different solver backends in parallel (greedy, metaheuristic, constraint programming, cloud APIs)
- Stress-test against 18 operational scenarios (cancellations, sick technicians, urgent injects, SLA shifts)
- Lock in-progress routes so re-optimization preserves real-world execution state
- Expose an Agentic AI layer that detects violations and triggers re-optimization autonomously
- Provide a benchmark studio with KPI charts, interactive maps, and Excel export
- Deploy to Hugging Face Spaces with pre-computed artifacts for instant demo

## Project Attributes

| Attribute         | Description                                                                                                                                                                               |
|-------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| problem-statement | Standard VRP solvers re-plan from scratch on every disruption — they're not aware that some stops are <i>already in execution</i> , leading to wasteful re-routing of stable assignments. |
| primary-objective | Build an execution-aware VRP platform that compares 8 solver families, locks in-progress work, and lets an AI agent autonomously fix constraint violations.                               |

| Attribute         | Description                                                                                                                                             |
|-------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------|
| core-technologies | Python 3.11, Streamlit, Google OR-Tools, PyVRP, OSRM, Esri ArcGIS, Google Gemini 2.0 Flash, Folium, Plotly, Docker, Hugging Face Spaces.                |
| repository-scope  | 8,140 lines across 4 source files · 18 fixed test scenarios · benchmark dataset · pre-computed artifacts · architecture docs · full function reference. |
| runtime-interface | Streamlit 6-tab UI: Overview · KPI Comparison · Map Explorer · Scenario Tables · Excel Export · 🤖 AI Agent.                                             |
| deployment-target | Hugging Face Spaces (Docker SDK) for public demo · Self-hosted Streamlit for production.                                                                |
| demo-surface      | Public HF Space with auto-loaded synthetic dataset and pre-computed benchmark artifacts.                                                                |
| key-capabilities  | 8 solver backends, locked-prefix re-optimization, agentic AI with 8 tools, interactive route maps, multi-sheet Excel export, reactive event handling.   |
| security-controls | Environment-driven API keys (GOOGLE_API_KEY) · Optional Esri portal authentication · No hardcoded credentials.                                          |
| production-focus  | Reproducible benchmarks · 100% assignment guarantees · zero-violation winner backend · sub-1000s runtime · cloud-deployable Docker image.               |

## Project Structure

```
src/content/posts/
├─ execution-aware-agentic-vrp/
│   ├── cover.png
│   └─ index.mdx

GitHub Repo: dranubhaparashar/Execution-Aware-Agentic-VRP-Solver-and-Benchmark-Studio/
├─ app.py                # Streamlit UI · 6 tabs · 845 lines
├─ base_engine.py        # 5 agents + scenario engine · 933 lines
├─ compare_backends.py   # 8 solver backends + helpers · 5,438 lines
├─ vrp_agent.py          # Gemini agent + 8 tools · 960 lines
├─ Dockerfile            # HF Spaces deployment
├─ requirements.txt
├─ data/                 # Auto-loaded on startup
│   ├── orders_40000.csv
│   ├── routes_400.csv
│   ├── depots_8.csv
│   ├── accounts.csv
│   └─ goal_profiles.csv
```

```
└─ artifacts/                                # Pre-computed benchmark results (8 backends)
|   └─ OR-Tools_Execution-Aware/
|       └─ vrp_scenario_results_full.xlsx
|   └─ OR-Tools/
|   └─ Hybrid_Execution-Aware_Rolling_VRP_Solver/
|   └─ Adaptive_Execution-Aware_Metaheuristic_Solver/
|   └─ Impacted-Subset_Greedy_2-opt/
|   └─ PyVRP/
|   └─ OSRM/
|       └─ Esri/
└─ docs/
    └─ ARCHITECTURE.md
    └─ ALGORITHM_COMPARISON.md
    └─ REFERENCE_*.md (4 function reference files)
```

## Why This Matters

**Note** A vehicle routing solver becomes operationally useful only when it is **execution-aware, comparable, and self-correcting**.

**Important** Real dispatch environments are not static — orders cancel, technicians call in sick, urgents inject mid-day. This project covers the missing layer between **batch route planning** and **live re-optimization with locked execution state**.

**Tip** The 8-backend benchmark format generalizes to any combinatorial routing problem — replace the scenarios and orders, the comparative scaffold remains useful.

**Warning** Lowest-distance solutions are not always best — PyVRP achieves the shortest routes (114 km) but produces 134,483 minutes of lateness. Constraint adherence matters as much as distance.

**Caution** Public demo deployments should never use production API keys, real customer data, or unrestricted solver access. Use the included synthetic dataset for showcases.

## What This System Does

flowchart TD

```
A[5 CSV files: orders, routes, depots, accounts, goal_profiles] -->
B[load_inputs_from_uploads]
B --> C[build_impacted_subset · 150 orders · 20 routes]
C --> D[ScenarioAgent applies 1 of 18 mutations]

D --> E{Backend Solver Selection}
E --> F1[OR-Tools EA · WINNER]
E --> F2[OR-Tools]
E --> F3[Hybrid]
E --> F4[Adaptive]
E --> F5[Greedy + 2-opt]
E --> F6[PyVRP]
E --> F7[OSRM]
E --> F8[Esri ArcGIS]

F1 & F2 & F3 & F4 & F5 & F6 & F7 & F8 --> G[_build_common_bundle]
G --> H[Output Bundle: 8 DataFrames]

H --> I[Streamlit UI · 6 tabs]
H --> J[Excel Export · 32 sheets]
H --> K[Interactive Folium Maps]
H --> L[VRPOrchestrator · Gemini Agent]

L --> M[8 Tools · Autonomous Re-optimization]
M --> H
```

## Agentic AI Request Flow

sequenceDiagram

autonumber

participant U as User / Dispatcher

participant ST as Streamlit UI

participant ORC as VRPOrchestrator

participant G as Gemini 2.0 Flash

participant T as Tool (Python fn)

participant CB as compare\_backends.py

U->>ST: "A driver called in sick – fix it"

ST->>ORC: handle\_event("RT\_6", reason)

ORC->>G: send\_message + 8 tool definitions

loop Up to 10 iterations

G-->>ORC: FunctionCall(detect\_violations)

ORC->>T: dispatch\_tool

T-->>ORC: Late routes + amounts

```
ORC->>G: FunctionResponse

G-->>ORC: FunctionCall(suggest_config)
ORC->>T: dispatch_tool
T-->>ORC: Config overrides + reasoning
ORC->>G: FunctionResponse

G-->>ORC: FunctionCall(trigger_reoptimization)
ORC->>CB: run_backend_suite(OR-Tools EA, new_config)
CB-->>ORC: New bundle
ORC->>G: Before/after diff

end

G-->>ORC: Final text summary
ORC-->>ST: "Re-routed 8 orders to neighbouring techs. Distance +3.2 km, lateness 0."
ST-->>U: Display in chat
```

## Core Capabilities

- 8 solver backends spanning greedy, metaheuristic, constraint programming, and cloud APIs
- 18 fixed operational scenarios for reproducible benchmarking
- Locked-prefix re-optimization that preserves in-progress technician work
- Agentic AI layer with 8 callable tools and 4 distinct agentic behaviors
- Interactive Folium route maps per scenario per backend
- 6-tab Streamlit UI with KPI charts, scenario tables, and Excel export
- Pre-computed artifacts that auto-load on app startup
- Hugging Face Spaces deployment via Docker
- BackendConfig with 28 tunable parameters exposed in the sidebar
- Reactive event buttons for sick driver, urgent order, cancellation, and shift change

## Bento Overview

| Layer                 | Description                                                                                     |
|-----------------------|-------------------------------------------------------------------------------------------------|
| Data Layer            | 40,000 orders · 400 routes · 8 depots · 4 goal profiles · auto-loaded from <code>./data/</code> |
| Scenario Layer        | 18 mutations across BASE, EAS_1-4, RT_1-7, GB_1-6 areas                                         |
| Solver Layer          | 8 backends with consistent input/output contract via <code>_build_common_bundle</code>          |
| Re-optimization Layer | Locked-prefix mechanism preserves stable assignments, only re-solves open work                  |

| Layer            | Description                                                                         |
|------------------|-------------------------------------------------------------------------------------|
| Agent Layer      | Gemini 2.0 Flash with function-calling, autonomous tool chaining, max 10 iterations |
| UI Layer         | Streamlit 6 tabs: Overview, KPI, Map, Scenario, Export, AI Agent                    |
| Artifact Layer   | Per-backend Excel + Folium HTML + PNG previews saved to ./artifacts/                |
| Deployment Layer | Docker image deployable to HF Spaces or any container host                          |

## The 8 Solver Backends — Benchmark Results

Tested on 150 orders · 20 routes · 18 scenarios from the included synthetic dataset.

| Rank | Backend                | Algorithm                       | Assigned   | BASE<br>km   | Late<br>min | OT<br>min | Runtime<br>s |
|------|------------------------|---------------------------------|------------|--------------|-------------|-----------|--------------|
| 🏆    | <b>OR-Tools<br/>EA</b> | Split-solve CP + GLS            | <b>150</b> | <b>231.8</b> | <b>0</b>    | <b>0</b>  | <b>770</b>   |
| 2    | OR-Tools               | CVRPTW + Guided<br>Local Search | 150        | 233.2        | 0           | 0         | 868          |
| 3    | Esri                   | Tabu Search (cloud)             | 142        | 238.7        | 34          | 0         | 973          |
| 4    | Adaptive               | Multi-restart regret            | 150        | 424.1        | 0           | 0         | 6,472        |
| 5    | Hybrid                 | Regret + LNS                    | 150        | 424.1        | 0           | 0         | 1,790        |
| 6    | Greedy                 | Nearest-neighbor +<br>2-opt     | 150        | 394.4        | 12,729      | 2,761     | 22           |
| 7    | PyVRP                  | Hybrid Genetic<br>Search        | 150        | 114.0        | 7,479       | 3,695     | 545          |
| 8    | OSRM                   | Greedy + road<br>distances      | 150        | 519.3        | 13,176      | 2,859     | 4,949        |

**Important OR-Tools EA is the only backend that achieves all four production requirements:** 100% assignment, zero violations, sub-1000s runtime, and locked-prefix support.

## The Innovation: Split-Solve Architecture

The OR-Tools Execution-Aware backend introduces a key architectural innovation. Locked orders **never enter the OR-Tools model** — they're copied verbatim from the baseline. Only open orders are submitted to the solver.

```
flowchart TD
    A["scenario_orders + scenario_routes + locked_prefixes"] --> B
    B["Step 1: Separate locked from open<br/>locked_order_ids = baseline<br/>assignments<br/>open_orders_df = orders NOT in locked set"]
    B --> C["Step 2: Anchor each vehicle<br/>last_locked.lat, last_locked.lon,<br/>last_locked.depart_time"]
    C --> D["Step 3: Build OR-Tools model with ONLY open orders<br/>Capacity = MaxOrderCount<br/>minus n_locked<br/>PATH_CHEAPEST_ARC + GUIDED_LOCAL_SEARCH<br/>Full 45s budget"]
    D --> E["Solution<br/>found?"]
    E -->|yes| F["Step 4: Merge output<br/>locked stops verbatim from baseline<br/>open stops<br/>from OR-Tools<br/>Recompute timing from anchor"]
    E -->|no or partial| G["Step 5: Greedy fallback<br/>Catches any dropped orders<br/>Guarantees<br/>150 of 150 assignment"]
    F --> H["Result: 150 of 150 assigned · 0 late · 0 overtime"]
    G --> H

    style H fill:#EAF3DE,stroke:#3B6D11,color:#27500A
    style D fill:#E6F1FB,stroke:#185FA5,color:#0C447C
```

### Why naive pinning fails

Submitting all 150 orders to OR-Tools with 130 of them having zero-slack time windows makes PATH\_CHEAPEST\_ARC infeasible. The solver fails to find any first solution within the time budget and returns 0 assignments.

By only submitting ~10 open orders, the model is **10-15x smaller**, the first solution comes instantly, and GLS gets the full 45s budget for improvement. Result: **11% faster than plain OR-Tools** and **0.6-3.4 km better** across all scenario groups.

---

## Detection and Solver Stack

### Solver Family Coverage



- **Constraint Programming:** OR-Tools, OR-Tools EA
- **Metaheuristics:** Hybrid, Adaptive, PyVRP
- **Greedy Heuristics:** Greedy + 2-opt, OSRM
- **Commercial Cloud:** Esri ArcGIS Network Analyst

## Scenario Coverage

- **Execution-Aware (EAS\_1-4):** Job timing changes mid-execution
- **Real-Time (RT\_1-7):** Cancellations, sick drivers, urgent injects, shift delays
- **Goal-Based (GB\_1-6):** Customer-specific weight overrides for travel/SLA/skill/overtime

## Output Bundle (per backend)

- scenario\_summary — 18 rows × 11 columns
- route\_output\_all — ~360 rows
- stop\_output\_all — ~2,700 rows
- order\_output\_all
- scenario\_actions
- scenario\_timings
- requirement\_checklist
- run\_meta

## Agentic AI Layer

```
import google.generativeai as genai
from google.generativeai import protos
import os, json

class VRPOrchestrator:
    def __init__(self, inputs, config, existing_results=None, progress_cb=None):
        self.inputs = inputs
        self.config = config
        self.results = existing_results or {}
        self._tool_calls_log = []
        self._chat_session = None

        genai.configure(api_key=os.environ["GOOGLE_API_KEY"])
        self._genai = genai
        self._model = genai.GenerativeModel(
            model_name="gemini-2.0-flash",
            tools=self._build_gemini_tools(),
            system_instruction=self._build_system_prompt(),
        )

    def chat(self, user_message: str) -> str:
```

```

        if self._chat_session is None:
            self._chat_session = self._model.start_chat(
                enable_automatic_function_calling=False
            )
        return self._run_agent_loop(user_message)

def _run_agent_loop(self, user_message: str) -> str:
    max_iterations = 10
    iteration = 0
    current_message = user_message

    while iteration < max_iterations:
        iteration += 1
        response = self._chat_session.send_message(current_message)

        fn_calls = []
        for candidate in response.candidates:
            for part in candidate.content.parts:
                if hasattr(part, "function_call") and part.function_call.name:
                    fn_calls.append(part.function_call)

        if not fn_calls:
            # Pure text response – agent is done
            return " ".join(p.text for c in response.candidates
                             for p in c.content.parts if hasattr(p, "text"))

        # Execute all function calls
        fn_responses = []
        for fn_call in fn_calls:
            result = self._dispatch_tool(fn_call.name, dict(fn_call.args))
            self._tool_calls_log.append({
                "tool": fn_call.name,
                "input": dict(fn_call.args),
                "result_preview": str(result)[:200],
            })
            fn_responses.append(
                self._genai.protos.Part(
                    function_response=self._genai.protos.FunctionResponse(
                        name=fn_call.name,
                        response={"result": json.dumps(result, default=str)},
                    )
                )
            )

        current_message = self._genai.protos.Content(
            role="user", parts=fn_responses,
        )

```

```
return "Agent reached max iterations."
```

## The 8 Agent Tools

| Tool                   | Purpose                                                   |
|------------------------|-----------------------------------------------------------|
| get_scenario_list      | Returns all 18 scenario IDs and descriptions              |
| explain_scenario       | Explains what a scenario tests + what mutation it applies |
| analyze_results        | Reads loaded bundles, surfaces metrics + violations       |
| detect_violations      | Finds late routes/orders + overtime amounts               |
| compare_backends       | Ranks backends by chosen priority criterion               |
| suggest_config         | Recommends BackendConfig overrides with reasoning         |
| run_backend            | Triggers a live solve and returns summary metrics         |
| trigger_reoptimization | Re-runs with improved config, returns before/after diff   |

## 4 Agentic Behaviors

1. **Natural language interpretation** — "Which backend is best for SLA?"
2. **Autonomous re-optimization** — agent detects violations, picks config, re-runs, compares
3. **Multi-step reasoning** — "Run A and B, then tell me which wins" chains 3+ tool calls
4. **Real-time reactive** — one-click event buttons trigger autonomous handling

## Dependency Stack

```
streamlit>=1.28,<2.0
pandas>=2.0
numpy>=1.24
ortools>=9.7
folium>=0.15
streamlit-folium>=0.15
plotly>=5.18
openpyxl>=3.1
requests>=2.31
google-generativeai>=0.5
nbformat>=5.9
matplotlib>=3.7
python-dateutil>=2.8
```

---

## Containerization Layer

```
FROM python:3.11-slim

RUN apt-get update && apt-get install -y --no-install-recommends \
    build-essential \
    && rm -rf /var/lib/apt/lists/*

RUN useradd -m -u 1000 user
USER user
ENV PATH="/home/user/.local/bin:$PATH"
WORKDIR /home/user/app

COPY --chown=user requirements.txt requirements.txt
RUN pip install --no-cache-dir --upgrade pip && \
    pip install --no-cache-dir -r requirements.txt

COPY --chown=user . .

EXPOSE 7860

CMD ["streamlit", "run", "app.py", \
    "--server.port=7860", \
    "--server.address=0.0.0.0", \
    "--server.headless=true", \
    "--server.enableCORS=false", \
    "--server.enableXsrfProtection=false", \
    "--browser.gatherUsageStats=false"]
```

---

## Deployment Path

```
flowchart LR
    A[Local Dev + Benchmark Excel] --> B[Extract per-backend artifacts]
    B --> C[Pre-populate ./artifacts/ folder]
    C --> D[Streamlit app + auto-load logic]
    D --> E[Docker build]
    E --> F[Push to Hugging Face Spaces]
    F --> G[Public HTTPS endpoint]
    G --> H[Auto-loaded artifacts on first view]
    G --> I[Live solve via run_backend_suite]
    G --> J[Agent chat with GOOGLE_API_KEY secret]
```

---

## Runtime Example

```
# Clone the repo
git clone https://github.com/dranubhaparashar/Execution-Aware-Agent-VRP-Solver-and-Benchmark-Studio.git
cd Execution-Aware-Agent-VRP-Solver-and-Benchmark-Studio

# Create virtual environment
python -m venv vrp
source vrp/bin/activate

# Install dependencies
pip install -r requirements.txt

# Set free Gemini API key from aistudio.google.com (no credit card)
export GOOGLE_API_KEY="YOUR_GOOGLE_API_KEY"

# Run the app
streamlit run app.py --server.port 8504

# Open http://localhost:8504
# Auto-loads synthetic data + pre-computed artifacts
# Click "Build comparison" to see all 8 backends instantly
# Open the AI Agent tab to chat with Gemini
```

## Hugging Face Deployment Pattern

- Build a Docker image with the Streamlit app and pre-computed artifacts
- Upload to Hugging Face Spaces (Docker SDK)
- Set GOOGLE\_API\_KEY as a Space secret
- Public HTTPS endpoint serves the auto-loaded studio
- Anyone can view benchmark results without running solvers
- Live re-solves available with one click per backend

```
git lfs install

git clone https://huggingface.co/spaces/AnubhaParashar/Execution-Aware-Agent-VRP-Solver-and-Benchmark-Studio
cd Execution-Aware-Agent-VRP-Solver-and-Benchmark-Studio

# Copy the entire vrp_hf_package contents into this folder
git lfs track "data/orders_40000.csv" # 7 MB file
git add .gitattributes
git add -A
git commit -m "Initial deploy"
git push
```

```
# Then in HF Settings → Variables and secrets:
```

```
#   GOOGLE_API_KEY = YOUR_GOOGLE_API_KEY
```

## Public Demo Surfaces

**Important** The Hugging Face Space provides instant browser-based access to the full studio with pre-loaded benchmark results. The GitHub repository contains the source code, architecture documentation, and reproducibility scripts.

## Demo Links

-  **Live App:** [AnubhaParashar / Execution-Aware-Agentic-VRP-Solver-and-Benchmark-Studio](#)
-  **GitHub Repository:** [Execution-Aware-Agentic-VRP-Solver-and-Benchmark-Studio](#)
-  **Wiki Documentation:** [Architecture · Algorithm Comparison · Technical Spec · Function Reference](#)

## Demo Video

**Embedded media:** [Open video](#)

## Algorithm Family Comparison

```
graph LR
    subgraph CP["Constraint Programming"]
        OR1[OR-Tools]
        OR2[OR-Tools EA · Winner]
    end

    subgraph META["Metaheuristics"]
        H[Hybrid]
        A[Adaptive]
        PV[PyVRP]
    end

    subgraph GR["Greedy Heuristics"]
        GD[Greedy + 2-opt]
        OS[OSRM]
    end

    subgraph CL["Commercial Cloud"]
        ES[Esri ArcGIS]
```

```
end
```

```
CP -->|Hard TW · 100 percent assign · Zero violations| WIN[Production-grade]  
META -->|Hard TW mostly · Slow · Custom rules| MIX[Specialized use]  
GR -->|TW in cost only · Fast · Violations| BASE[Baseline only]  
CL -->|Soft TW · Drops orders · Real roads| ENT[Enterprise GIS]
```

```
style OR2 fill:#EAF3DE,stroke:#3B6D11,color:#27500A
```

---

## Engineering Value

This project is not just about routing optimization.

It demonstrates how to take a combinatorial optimization problem and make it:

- reproducible (18 fixed scenarios + synthetic dataset)
- comparable (8 backends with identical I/O contract)
- execution-aware (locked-prefix preserves real-world state)
- agentic (Gemini autonomously detects and fixes violations)
- containerized (Docker image deployable to any host)
- publicly demoable (Hugging Face Space with pre-loaded artifacts)
- documented for reuse (Architecture, Algorithm Comparison, Function Reference)

---

## Current Strengths

- Single Python codebase covers all 8 solver families uniformly
- Locked-prefix mechanism preserves in-progress assignments during re-optimization
- OR-Tools EA achieves 11% faster runtime + better distance than plain OR-Tools
- Agentic AI layer makes the system self-correcting without human dispatch input
- Pre-computed artifacts allow instant demo without running solvers
- Free Gemini API tier (1,500 req/day) keeps the AI layer zero-cost
- Reactive event buttons cover the 4 most common operational disruptions
- Excel export captures all 32 sheets for offline analysis
- Folium maps render route visualizations interactively per backend per scenario

---

## Next Improvements

- Add CI/CD pipeline for automated benchmark regression on new scenarios
- Persist OSRM cache to disk to survive Space restarts
- Add multi-depot solving support
- Add real-time technician GPS feed integration
- Add SLA breach prediction model alongside the routing solver

- Add A/B testing framework for cost function variants
  - Add OpenAPI spec for headless solver invocation
  - Move to Anthropic Claude as alternative agent backend
  - Add Kubernetes deployment manifest for production scale-out
  - Add load testing with realistic concurrent user patterns
- 

## Key Innovation

**Important** This project connects **vehicle routing optimization** with **execution-aware re-planning** and **agentic AI orchestration** — three layers that are usually built separately.

It turns:

**Static Route Planning → Execution-Aware Re-Optimization → Autonomous Violation Detection and Recovery**

rather than stopping at one-shot solver output.

The OR-Tools Execution-Aware split-solve architecture is the keystone — it shows that **operational state (which stops are committed) is just as important as the optimization objective**.

---

## Conclusion

This repository shows how a combinatorial optimization problem can evolve from a single solver invocation into a comparative benchmark studio with autonomous AI orchestration.

It combines:

- 8 solver backends with consistent I/O contract
  - 18 reproducible operational scenarios
  - Locked-prefix execution-awareness
  - Google Gemini agentic AI with 8 callable tools
  - Streamlit interactive UI with 6 specialized tabs
  - Pre-computed artifacts for instant demo
  - Docker deployment to Hugging Face Spaces
  - Comprehensive technical documentation (Architecture, Algorithm Comparison, Function Reference)
- 

## Final Thought



From **solving a routing problem** to **shipping an autonomous, execution-aware optimization platform**

The real value is not only the OR-Tools EA winner backend — it is the **comparative benchmark scaffolding, the locked-prefix mechanism, and the agentic AI layer** that together make the system operationally trustworthy.

100% assignment · 0 violations · 770s runtime · 8 solvers compared · 18 scenarios stress-tested  
· 1 AI agent ready to handle the next disruption.

Source: <src/content/posts/execution-aware-agentic-vrp/execution-aware-agentic-vrp.md>