

Self Project

2026-06-23

No cover image available.

LightDID-ZKP: Policy- and Resource-Aware Selection of BBS and AnonCreds Verifiable Presentations

LightDID-ZKP is a self research project that introduces CAPS-ZK, a policy- and resource-aware selector for choosing between BBS selective-disclosure presentations and AnonCreds predicate presentations in privacy-preserving decentralized identity systems.

Project Snapshot

DATE

2026-06-23

CATEGORY

Self Project

SHORT DESCRIPTION

LightDID-ZKP is a self research project that introduces CAPS-ZK, a policy- and resource-aware selector for choosing between BBS selective-disclosure presentations and AnonCreds predicate presentations in privacy-preserving decentralized identity systems.

TAGS / TECH STACK

Decentralized Identity

Verifiable Credentials

Zero-Knowledge Proofs

BBS

AnonCreds

Selective Disclosure

Privacy-Preserving Identity

CAPS-ZK

Python

Research Prototype

Project Links

[GitHub Link](#)[GitHub Link](#)[GitHub Link](#)[GitHub Link](#)[GitHub Link](#)[GitHub Link](#)[GitHub Link](#)[GitHub Link](#)[GitHub Link](#)[GitHub Link](#)[GitHub Link](#)[GitHub Link](#)[GitHub Link](#)[GitHub Repository](#)

Full Project Content

LightDID-ZKP is a policy- and resource-aware framework for selecting the correct privacy-preserving verifiable presentation backend in decentralized identity systems. It focuses on choosing between BBS selective-disclosure presentations and AnonCreds predicate presentations without allowing unsafe privacy downgrades or cost-only fallback decisions.

GitHub Repository: [dranubhaparashar/LightDID-ZKP](#)

Wiki documentation: [Home](#) · [Architecture](#) · [CAPS-ZK Selection Logic](#) · [Policy Model](#) · [Experiment Protocol](#) · [Results and Figures](#)

Reproducibility: [Reproducing the Paper Tables](#) · [Optional Real-Backend Benchmarks](#)

Claim boundary: [Security, Privacy, and Claim Boundary](#).

One-Line Idea

LightDID-ZKP selects the right privacy-preserving presentation mechanism by checking verifier policy semantics, privacy requirements, credential-backed predicate needs, and device resource limits before applying cost scoring.

The core idea is simple:

Do not choose the cheapest proof first. Choose only from mechanisms that satisfy the verifier policy and privacy semantics, then optimize for cost.

Why This Project Exists

Decentralized identity systems increasingly support privacy-preserving verifiable presentations. A holder wallet may support more than one presentation mechanism, but each mechanism is suitable for a different kind of verifier request.

For example, a verifier may ask for selected attributes such as name, degree, institution, or membership status. In such cases, a BBS-style selective-disclosure presentation can reveal only the required fields while hiding the rest.

Another verifier may ask for a predicate such as age greater than 18, income above a threshold, eligibility status, or another credential-backed condition. In that case, simple selective disclosure is not enough because the verifier needs proof that a hidden predicate is satisfied.

LightDID-ZKP addresses this decision problem. It introduces a selector that does not treat all proof systems as interchangeable. Instead, it checks the policy, privacy level, predicate requirement, issuer binding, freshness rule, status-policy requirement, and device constraints before selecting the presentation backend.

Project at a Glance

Area	Description
Full Form	Lightweight Decentralized Identity Zero-Knowledge Presentation Selector
Project Type	Self research project and reproducible prototype
Primary Goal	Select between BBS and AnonCreds verifiable presentation backends using policy and resource constraints
Core Selector	CAPS-ZK: Capability- and Policy-Safe Zero-Knowledge selector
Main Mechanisms	BBS selective disclosure and AnonCreds predicate presentation
Main Users	DID wallet developers, identity researchers, privacy engineers, academic reviewers, ZKP learners
Input	Verifier policy, holder capabilities, credential metadata, privacy requirement, resource profile
Output	Selected presentation mechanism, rejection reason, cost score, policy-safety explanation
Main Repository	dranubhaparashar/LightDID-ZKP
Documentation Surface	GitHub README and Wiki pages
Claim Boundary	Research prototype, not a production DID wallet or new cryptographic primitive

What Makes LightDID-ZKP Different

Important

LightDID-ZKP is not a new cryptographic proof primitive. It is a selector and orchestration layer above existing presentation mechanisms.

Ordinary Cost-First Choice	LightDID-ZKP Approach
Chooses the lowest latency backend	First checks whether the backend satisfies verifier policy semantics
Treats disclosure and predicate proof as interchangeable	Distinguishes selective disclosure from credential-backed predicates
Falls back to a cheaper proof even when semantics differ	Rejects unsafe fallback
Optimizes only for size or speed	Balances policy fit, privacy level, latency, VP size, verification cost, and memory
Gives weak explanation for backend choice	Produces a reasoned selector decision
Ignores claim boundary	Clearly states what the prototype does and does not prove

Reader Walkthrough

A reader can understand LightDID-ZKP in five steps:

1. **Verifier sends a presentation request** containing required attributes, predicate needs, issuer requirements, freshness expectations, status-policy rules, and resource constraints.
2. **CAPS-ZK parses the policy** and identifies whether the request requires selective disclosure, credential-backed predicates, or both.
3. **Candidate mechanisms are filtered** using semantic, privacy, predicate, issuer-binding, freshness, status-policy, and resource checks.
4. **Only safe mechanisms are scored** using latency, proof size, verification cost, and memory indicators.
5. **The selected backend is returned** with a reason. If no backend satisfies the policy safely, the selector rejects the request instead of downgrading privacy or semantics.

System Architecture

flowchart LR

```
I[Issuer] --> C[Verifiable Credential]
C --> W[Holder Wallet]

W --> P[Verifier Policy Parser]
P --> S[CAPS-ZK Selector]

S --> BBS[BBS Selective Disclosure]
S --> AC[AnonCreds Predicate Presentation]

BBS --> VP[Verifiable Presentation]
AC --> VP

VP --> V[Verifier]
V --> R[Policy-Compliant Decision]
```

LightDID-ZKP separates credential issuance, holder-side decision logic, backend selection, cryptographic presentation generation, and verifier-side policy checking.

CAPS-ZK does not replace BBS or AnonCreds. It decides when each mechanism is appropriate and prevents unsafe substitutions.

CAPS-ZK Selection Workflow

flowchart TD

```
A[Presentation Request] --> B[Read Verifier Policy]
B --> C{Credential-backed predicate required?}

C -->|Yes| D[Check AnonCreds support]
C -->|No| E[Check BBS selective disclosure]

D --> F{Policy semantics satisfied?}
E --> F

F -->|No| G[Reject unsafe mechanism]
F -->|Yes| H[Check privacy and issuer constraints]

H --> I[Check freshness and status policy]
I --> J[Check holder resource profile]
J --> K[Score latency, VP size, verification cost, memory]
K --> L[Select policy-safe mechanism]
L --> M[Return decision and reason]
```

The workflow is intentionally policy-first. Resource optimization is applied only after the mechanism passes the safety and semantics checks.

Core Modules

1. Policy Model

The policy model represents what the verifier is actually asking for. It includes attribute disclosure requirements, predicate requirements, issuer binding, freshness rules, credential status rules, privacy level, and domain binding.

The goal is to make the selector understand the difference between:

- revealing a selected attribute, and
- proving a hidden predicate about an attribute.

This distinction is important because selective disclosure should not be used as an unsafe substitute for a required hidden predicate.

2. Capability Model

The capability layer describes what each backend can support. BBS is modeled for selective disclosure. AnonCreds is modeled for credential-backed predicates. CAPS-ZK checks these capabilities before using a backend.

3. Privacy-Level Abstraction

The selector assigns a privacy meaning to each mechanism. A presentation that reveals an attribute is not privacy-equivalent to a predicate proof that keeps the attribute hidden.

This allows LightDID-ZKP to prevent privacy downgrades.

4. Resource Profiler

The resource profile captures holder-side constraints such as latency budget, available memory, verification cost tolerance, and maximum acceptable verifiable presentation size.

This is useful for wallets running on constrained devices or mobile-like environments.

5. Cost Scorer

After unsafe mechanisms are removed, the remaining candidates are scored using measured or configured cost indicators. The cost scorer helps pick the most practical valid backend.

6. Backend Adapter Layer

The backend adapter layer keeps the selector independent from specific cryptographic implementation details. The repository includes optional real-backend templates so the framework can be extended with concrete BBS or AnonCreds libraries.

Presentation Profiles

Profile	Main Purpose	Typical Selection Case
BBS selective disclosure	Reveal selected attributes from a signed credential	Attribute disclosure without hidden predicate proof
AnonCreds predicate presentation	Prove credential-backed predicates	Age, threshold, eligibility, or range-style requirements
CAPS-ZK selector	Select the correct presentation backend	Policy-aware and resource-aware wallet decision
Policy-safe rejection	Refuse unsafe fallback	No backend satisfies required semantics safely

End-to-End Workflow

Stage	What Happens	Output
1. Presentation Request	Verifier requests disclosure or predicate proof	Policy object
2. Policy Parsing	Selector extracts semantic and privacy requirements	Normalized verifier policy
3. Capability Filtering	Backends that cannot satisfy the request are removed	Valid candidate set
4. Privacy Filtering	Mechanisms that downgrade privacy are rejected	Privacy-safe candidate set
5. Issuer and Freshness Checks	Issuer, status, and freshness rules are validated	Policy-compliant candidate set
6. Resource Filtering	Device limits are checked	Runtime-feasible candidates
7. Cost Scoring	Latency, VP size, verification time, and memory are compared	Ranked backend list
8. Selection or Rejection	Best safe backend is selected, or request is rejected	Decision and reason
9. Presentation Generation	Chosen backend generates the VP	Verifiable presentation

Stage	What Happens	Output
10. Verifier Decision	Verifier checks the VP and policy compliance	Accept/reject result

Example Decision Cases

Case A: Selective Attribute Disclosure

A verifier asks the holder to reveal only:

- name,
- degree,
- institution.

No hidden predicate is required.

Expected selector behavior: choose BBS selective disclosure if it satisfies the issuer, freshness, status, and resource constraints.

Case B: Age Predicate

A verifier asks the holder to prove:

- age is greater than or equal to 18,
- without revealing the exact date of birth.

Expected selector behavior: choose AnonCreds predicate presentation if credential-backed predicate proof is required.

Case C: Unsafe Cost-First Fallback

A cheaper backend exists, but it cannot prove the required predicate.

Expected selector behavior: reject unsafe fallback. LightDID-ZKP should not replace a hidden predicate proof with simple disclosure just because it is cheaper.

Experiment Protocol

The experiment package follows a manuscript-style reproducibility setup.

Item	Configuration
Warmup runs	5
Measured runs	50

Item	Configuration
Attribute counts	4, 8, 16, 32, 64
Mechanisms	BBS and AnonCreds
Main metrics	Proving latency, verification latency, VP size, RSS memory, coefficient of variation
Supporting studies	Selector decisions, ablation, resource sensitivity
Outputs	Tables, figures, CSV summaries, decision logs

The repository contains benchmark summaries, selector decisions, ablation results, resource-sensitivity outputs, and generated figures.

Repository Structure

```
LightDID-ZKP
├─ assets/
│   └─ diagrams/
│       ├── lightdid_banner.svg
│       ├── lightdid_layered_architecture.svg
│       ├── caps_zk_selection_flow.svg
│       ├── experiment_pipeline.svg
│       └─ verifier_metadata_guard.svg
│
├─ benchmarks/
│   ├── lightdid_benchmark_summary.csv
│   ├── selector_decisions.csv
│   ├── resource_sensitivity.csv
│   └─ ablation_cost_first_fallback.csv
│
├─ configs/
│   ├── policies.yaml
│   ├── device_profiles.yaml
│   └─ experiment_config.yaml
│
├─ docs/
│   ├── ARCHITECTURE.md
│   ├── DEVELOPER_GUIDE.md
│   ├── REVIEWER_NOTES.md
│   └─ GITHUB_ABOUT.md
│
├─ experiments/
│   ├── run_all.py
│   └─ generate_tables.py
```

```

|   ├── plot_results.py
|   └── optional_real_backend_templates/
|
| ├── results/
|   ├── figures/
|   └── tables/
|
| ├── src/
|   ├── lightdid_zkp/
|   │   ├── selector.py
|   │   ├── policy.py
|   │   ├── profiles.py
|   │   ├── metrics.py
|   │   └── utils.py
|   └──
|
| ├── tests/
|   └── test_selector.py
|
| ├── README.md
| ├── requirements.txt
| ├── pyproject.toml
| ├── CITATION.cff
| └── LICENSE

```

Important Project Paths

Component	Path
Main selector code	src/lightdid_zkp/selector.py
Policy model	src/lightdid_zkp/policy.py
Device profiles	src/lightdid_zkp/profiles.py
Experiment runner	experiments/run_all.py
Table generation	experiments/generate_tables.py
Plot generation	experiments/plot_results.py
Policy configs	configs/policies.yaml
Device configs	configs/device_profiles.yaml
Benchmark CSV files	benchmarks/
Result tables	results/tables/

Component	Path
Result figures	results/figures/
Architecture diagrams	assets/diagrams/
Developer notes	docs/
Tests	tests/

Local Run Command

Create and activate a Python environment:

```
python -m venv .venv
```

Windows:

```
.venv\Scripts\activate
```

macOS/Linux:

```
source .venv/bin/activate
```

Install dependencies:

```
pip install -r requirements.txt
```

Run the full experiment pipeline:

```
python experiments/run_all.py
```

Run tests:

```
pytest -q
```

Reproduce Tables and Figures

Generate manuscript-style tables:

```
python experiments/generate_tables.py
```

Generate result figures:

```
python experiments/plot_results.py
```

Outputs are saved in:

```
results/tables/  
results/figures/
```

Example Selector Usage

```
from lightdid_zkp.selector import select_presentation  
from lightdid_zkp.policy import VerifierPolicy  
from lightdid_zkp.profiles import DeviceProfile  
  
policy = VerifierPolicy(  
    required_attributes=["name", "degree", "institution"],  
    predicate_requirements=[],  
    max_vp_size_kb=32,  
    require_credential_backed_predicate=False,  
)  
  
device = DeviceProfile(  
    name="mobile_like",  
    max_latency_ms=500,  
    max_memory_mb=512,  
)  
  
decision = select_presentation(policy=policy, device=device)  
  
print("Selected mechanism:", decision.mechanism)  
print("Reason:", decision.reason)
```

Example output for selective disclosure:

```
Selected mechanism: BBS  
Reason: Policy requires selective disclosure only; BBS satisfies disclosure constraints with  
lower estimated presentation size.
```

Example output for predicate-heavy policies:

Selected mechanism: AnonCreds

Reason: Verifier policy requires credential-backed predicate proof; AnonCreds satisfies predicate semantics.

Project Links

GitHub Repository: [dranubhaparashar/LightDID-ZKP](#)

- **Wiki Home:** [LightDID-ZKP Wiki](#)
- **Architecture:** [Architecture and Design](#)
- **CAPS-ZK Logic:** [CAPS-ZK Selection Logic](#)
- **Policy Model:** [Policy Model](#)
- **Experiments:** [Experiment Protocol](#)
- **Results:** [Results and Figures](#)
- **Claim Boundary:** [Security, Privacy, and Claim Boundary](#)

Documentation and Wiki Links

Because this project does not currently have a YouTube walkthrough, the Wiki is the main public explanation surface.

Wiki Page	Purpose
Home	Main entry point for the project
About LightDID-ZKP	Short overview and motivation
Architecture and Design	System design and component view
CAPS-ZK Selection Logic	Selector logic and filtering sequence
Policy Model	Formal policy inputs and constraints
Experiment Protocol	Benchmarking and evaluation setup
Results and Figures	Output tables and generated figures
Reproducing the Paper Tables	Steps for reproducibility
Optional Real-Backend Benchmarks	Notes for real BBS/AnonCreds backend experiments
Repository Structure	File and folder explanation
Security, Privacy, and Claim Boundary	What the project does and does not claim

Wiki Page	Purpose
Developer Guide	Developer notes and extension points
Quick Links	Fast navigation page

Architecture Diagrams

The repository includes project diagrams under `assets/diagrams/`.

Diagram	Link
LightDID-ZKP Banner	<code>assets/diagrams/lightdid_banner.svg</code>
Layered Architecture	<code>assets/diagrams/lightdid_layered_architecture.svg</code>
CAPS-ZK Selection Flow	<code>assets/diagrams/caps_zk_selection_flow.svg</code>
Experiment Pipeline	<code>assets/diagrams/experiment_pipeline.svg</code>
Verifier Metadata Guard	<code>assets/diagrams/verifier_metadata_guard.svg</code>

Result Files

The repository includes result files for:

- BBS presentation latency,
- AnonCreds presentation latency,
- verification time,
- verifiable presentation size,
- RSS memory usage,
- coefficient of variation,
- selector decisions across policy types,
- ablation comparison against unsafe cost-first fallback,
- resource-sensitivity analysis.

Result folders:

```
benchmarks/
results/tables/
results/figures/
```

Research Scope and Claim Boundary

LightDID-ZKP is intended as a research prototype for privacy-preserving decentralized identity presentation selection.

It does **not** claim to:

- introduce a new zero-knowledge proof primitive,
- replace BBS or AnonCreds,
- provide production wallet security certification,
- prove constrained-device deployment without real device measurements,
- implement a complete DID wallet,
- provide legal or compliance certification for digital identity deployments.

It does claim to:

- provide a policy-aware selection framework,
- provide a resource-aware presentation decision layer,
- compare presentation-level trade-offs,
- support reproducibility through scripts, tables, and figures,
- demonstrate why policy semantics must be checked before cost optimization.

Practical Use Cases

User Group	How LightDID-ZKP Helps
DID wallet developers	Provides a policy-safe way to choose between multiple presentation backends
Identity researchers	Demonstrates a formal selector model for privacy-preserving verifiable presentations
ZKP learners	Shows why selective disclosure and predicate proofs are not interchangeable
Academic reviewers	Provides reproducible scripts, tables, diagrams, and clear claim boundaries
Privacy engineers	Highlights unsafe fallback and privacy-downgrade risks
Standards readers	Connects verifier policy semantics with presentation backend selection
Project evaluators	Provides GitHub repository, Wiki documentation, diagrams, and experiment outputs

Engineering Value

LightDID-ZKP demonstrates how decentralized identity presentation selection can be treated as an engineering decision problem rather than only a cryptographic primitive choice.

The project connects:

- verifier policy parsing,
- backend capability modeling,
- privacy-level abstraction,
- predicate requirement checking,
- resource filtering,
- benchmark-driven cost scoring,
- safe rejection,
- reproducible experiment generation,
- GitHub documentation and Wiki publishing.

This makes the project useful as a research prototype, academic support repository, and technical portfolio project.

Current Strengths

- Clear focus on decentralized identity and verifiable presentation selection.
 - Separates policy correctness from cost optimization.
 - Distinguishes BBS selective disclosure from AnonCreds predicate presentation.
 - Prevents unsafe fallback and privacy downgrades.
 - Includes reproducibility scripts, tables, result files, and diagrams.
 - Provides GitHub Wiki documentation instead of relying only on README text.
 - States a clear claim boundary, which is important for research credibility.
 - Suitable for manuscript support, GitHub portfolio display, and technical review.
-

Key Innovation

Key point

The strongest innovation is not claiming a new proof system. The strongest innovation is the policy-safe selector: LightDID-ZKP checks whether a presentation backend satisfies verifier semantics and privacy requirements before optimizing for resource cost.

Most systems discuss proof mechanisms individually. LightDID-ZKP focuses on the wallet-side decision problem that appears when multiple privacy-preserving presentation mechanisms are available.

Conclusion

LightDID-ZKP shows how decentralized identity wallets can make safer presentation choices when multiple privacy-preserving mechanisms are available.

It connects verifier policy, privacy semantics, backend capability, resource feasibility, and cost scoring into one reproducible framework.

The project is valuable because it explains a practical identity-system problem: a low-cost proof is not acceptable if it does not satisfy the verifier's actual requirement. By placing policy safety before cost optimization, LightDID-ZKP provides a clean and credible research direction for privacy-preserving decentralized identity presentation selection.

Final Thought

Policy first. Privacy second. Cost only after safety.

LightDID-ZKP is about making decentralized identity presentation selection safer, explainable, and reproducible.

BBS selective disclosure · AnonCreds predicates · CAPS-ZK selector · no unsafe fallback · no privacy downgrade · GitHub Wiki documentation

Source: [src/content/posts/LightDID-ZKP- A policy and resource aware framework for selecting BBS and AnonCreds verifiable presentations in decentralized identity/index.md](#)