


Industrial Project

2026-05-02

Predictive & Preventive Maintenance for Generator Reliability

A full-stack analytics and ML dashboard that measures PM effectiveness, forecasts 14/30-day generator failure risk, and recommends prioritized maintenance actions.



Project Snapshot

DATE

2026-05-02

CATEGORY

Industrial Project

SHORT DESCRIPTION

A full-stack analytics and ML dashboard that measures PM effectiveness, forecasts 14/30-day generator failure risk, and recommends prioritized maintenance actions.

TAGS / TECH STACK

AI

Predictive Maintenance

Machine Learning

Streamlit

Scikit-learn

Generator Reliability

MLOps

Hugging Face

Snowflake

Python

Project Links

Demo Video

Hugging Face Space

GitHub Link

Project Link

Demo Video

GitHub Link

GitHub Link

Full Project Content

Building a maintenance analytics platform that doesn't just **track failures** — but links PM events to subsequent breakdowns, trains classification models on telemetry signals, and turns risk scores into prescriptive maintenance plans.

GitHub Repository:

[dranubhaparashar/Predictive Preventive Maintenance for Generator Reliability](https://github.com/dranubhaparashar/Predictive-Preventive-Maintenance-for-Generator-Reliability)

 **Live demo video:** youtube.com/watch?v=QQSSFwxY_ro

 **Try it live:** [Hugging Face Space](#)

Vision

Most maintenance teams know a generator failed. Very few can answer the harder question:

After preventive maintenance is completed, how long does the generator actually run before the next failure?

This project builds the full analytical chain that answers that question — and goes further, training ML models to predict which assets are about to fail next, then combining both into a prioritized action plan:

- Link every PM event to its subsequent failure for the same asset
- Compute `days_to_next_failure` as a measurable PM effectiveness metric
- Train Logistic Regression and Random Forest classifiers on weekly telemetry signals
- Predict `failure_within_14d` and `failure_within_30d` per asset
- Synthesize historical patterns, ML scores, and business-impact data into maintenance recommendations
- Expose everything through a 9-tab Streamlit dashboard with CSV upload, KPI summaries, and downloadable outputs
- Deploy to Hugging Face Spaces (Docker), Snowflake Streamlit, and local Python environments

Project Attributes

Attribute	Description
problem-statement	Maintenance teams treat PM as a cost with no measurable return. Without linking PM events to subsequent failures, there is no way to know whether scheduled maintenance is improving reliability or being done at the wrong intervals.

Attribute	Description
primary-objective	Build a data-driven platform that quantifies PM effectiveness, scores asset-level failure probability, and recommends prioritized maintenance actions.
core-technologies	Python, Streamlit, Pandas, NumPy, Scikit-learn, Plotly, Matplotlib, Docker, Hugging Face Spaces, Snowflake Streamlit.
dataset-structure	7 CSV files: asset master, PM events, failure events, linked PM-failure pairs, weekly telemetry, labeled telemetry for ML, business impact.
runtime-interface	Streamlit 9-tab UI: Problem · Dataset · Executive KPIs · PM Effectiveness · Failures & Cost · Telemetry Risk · Predictive ML & PM Strategy · Saved Outputs · Data Explorer.
deployment-target	Hugging Face Spaces (Docker SDK) · Snowflake Streamlit (Python 3.11) · Local Python.
demo-surface	Public HF Space with embedded synthetic dataset; live model training in-browser.
key-capabilities	PM-to-failure event linking, on-time vs. delayed PM comparison, 14/30-day failure classification, feature importance, asset risk ranking, prescriptive action planner, KPI exports.
production-focus	Reproducible on real PM work orders and failure tickets with retraining; synthetic demo ships with the repo.

Three-Phase Architecture

```
flowchart TD
    A["A[7 CSV files: asset_master, pm_events, failure_events,\npm_failure_linked, telemetry_weekly,\ngenerator_telemetry_with_labels, business_impact]"] --> B["B[Phase 1 · Reliability Analytics]"]
    B --> C["C[Link PM → next failure per asset_id\nCompute days_to_next_failure]"]
    C --> D["D[Phase 2 · Predictive ML]"]
    D --> E["E[Feature engineering on telemetry_weekly\nLogistic Regression baseline\nRandom Forest advanced]"]
    E --> F["F[failure_within_14d\nfailure_within_30d\nper-asset probability scores]"]
    F --> G["G[Phase 3 · Prescriptive Strategy]"]
    G --> H["H[Merge analytics + ML risk + PM history + business impact]"]
    H --> I["I[Action: Immediate Inspection\nPM within 3-7 days\nMonitoring\nNormal Operations]"]
    I --> J["J[9-tab Streamlit Dashboard]"]
    J --> K["K[KPI charts · Telemetry anomaly viz · Risk ranking · Excel export]"]
```

Why This Matters

Note A generator failure during a critical operation is not just a repair cost — it is downtime, SLA breach, customer revenue loss, and safety exposure. The difference between reactive and predictive maintenance is the ability to act **before** the failure event.

Important PM effectiveness cannot be assumed. On-time PM may still result in early failure if intervals are wrong, parts degrade faster in certain environments, or specific asset models have higher intrinsic risk. Only by linking PM events to subsequent failures can teams validate whether their schedule is working.

Tip The `days_to_next_failure` metric is deceptively simple. It compresses the entire PM-effectiveness question into a single comparable number per PM event, enabling distribution analysis, regional comparison, and model-type segmentation without requiring complex survival models.

Warning ML predictions are only as good as the failure labels. The synthetic dataset ships with realistic label distributions, but production deployment requires genuine labeled failure outcomes and scheduled model retraining as operational patterns shift.

Caution Business-impact figures (revenue loss, SLA penalties, customer exposure) in the demo are estimates based on synthetic data. Production use requires connecting real financial and SLA records to the failure event table.

Dataset Structure

File	Purpose
<code>asset_master.csv</code>	Generator inventory: <code>asset_id</code> , <code>model</code> , <code>region</code> , <code>criticality</code> , <code>environment_type</code> , <code>install_date</code>
<code>pm_events.csv</code>	Scheduled and completed PM records with dates, technician, completion status, cost

File	Purpose
failure_events.csv	Repair tickets with downtime hours, severity, root cause category, repair cost
pm_failure_linked.csv	Derived: each PM event linked to the next failure for that asset, with days_to_next_failure
telemetry_weekly.csv	Weekly sensor readings: temperature, vibration, fuel rate, alarm count, anomaly_score
generator_telemetry_with_labels.csv	Telemetry enriched with failure_within_14d and failure_within_30d binary labels
business_impact.csv	Revenue loss per outage, SLA breach counts, affected customers, financial exposure

Core Linking Algorithm

For each PM event, the platform:

1. Identifies the `asset_id` and PM completion date
2. Queries all failure events for that asset after the PM date
3. Selects the earliest subsequent failure
4. Computes `days_to_next_failure = failure_date - pm_completion_date`

This produces one row per PM event with a measurable survival duration — the foundation for all downstream analytics and model training.

```
def link_pm_to_next_failure(pm_events: pd.DataFrame, failure_events: pd.DataFrame) ->
pd.DataFrame:
    records = []
    for _, pm in pm_events.iterrows():
        asset_failures = failure_events[
            (failure_events["asset_id"] == pm["asset_id"]) &
            (failure_events["failure_date"] > pm["pm_completion_date"])
        ].sort_values("failure_date")

        if not asset_failures.empty:
            next_failure = asset_failures.iloc[0]
            records.append({
                "pm_id": pm["pm_id"],
                "asset_id": pm["asset_id"],
                "pm_completion_date": pm["pm_completion_date"],
                "next_failure_date": next_failure["failure_date"],
                "days_to_next_failure": (
```

```
        next_failure["failure_date"] - pm["pm_completion_date"]
    ).days,
    "pm_on_time": pm["pm_on_time"],
    "failure_severity": next_failure["severity"],
    })
return pd.DataFrame(records)
```

Machine Learning Pipeline

Feature Set

Feature	Source	Why It Matters
days_since_last_pm	pm_events	Longer gaps increase failure risk
runtime_hours_week	telemetry_weekly	High utilization accelerates wear
oil_temperature	telemetry_weekly	Thermal stress indicator
coolant_temperature	telemetry_weekly	Cooling system degradation signal
battery_voltage	telemetry_weekly	Electrical system health
vibration	telemetry_weekly	Mechanical wear proxy
fuel_rate	telemetry_weekly	Efficiency drop = early warning
alarm_count	telemetry_weekly	Cumulative fault indicators
anomaly_score	telemetry_weekly	Composite deviation metric
model, region, criticality, environment_type	asset_master	Asset-class risk stratification

Models

Model	Role	Strength
Logistic Regression	Baseline	Interpretable coefficients, fast inference
Random Forest	Production	Captures nonlinear interactions, feature importance ranking

Outputs per Model

- Accuracy, Precision, Recall, F1, ROC AUC

- Confusion matrices (14d and 30d horizons)
- Feature importance bar charts
- Per-asset failure probability scores
- Risk band classification: Critical / High / Medium / Low

Prescriptive Action Logic

```
flowchart TD
    A[Asset risk score from Random Forest] --> B{Risk band}
    B -->|Critical · P > 0.75| C[Immediate Inspection]
    B -->|High · P 0.50-0.75| D[PM within 3-7 days]
    B -->|Medium · P 0.25-0.50| E[Monitoring · next scheduled PM]
    B -->|Low · P < 0.25| F[Normal Operations]

    C --> G[Cross-check: days_since_last_pm > threshold?]
    G -->|Yes| H[Escalate: Emergency Work Order]
    G -->|No| I[Standard Inspection Ticket]

    D --> J[Check business_impact: revenue_at_risk]
    J -->|High revenue exposure| K[Priority queue]
    J -->|Normal| L[Standard queue]
```

Dashboard Tabs

Tab	Content
Problem & Solution	Business context: the cost of reactive maintenance and the PM-effectiveness gap
Dataset & Raw → Algorithm	CSV preview, PM linking logic, enrichment walkthrough
Executive KPIs	Total downtime, repair costs, PM completion rate, SLA breach count
PM Effectiveness	days_to_next_failure distribution, on-time vs. delayed PM comparison, model/region breakdown
Failures & Cost	Failure severity breakdown, root cause categories, cost per failure, top assets by downtime
Telemetry Risk	Anomaly score trends, sensor heatmaps, aging asset identification
Predictive ML & PM Strategy	Model training, confusion matrices, feature importance, risk rankings, action plan table

Tab	Content
Saved Outputs	Exported figures and result tables from current session
Data Explorer & Export	Filtered dataset views, downloadable CSV and Excel exports

Bento Overview

Layer	Description
Data Layer	7 CSV files covering asset inventory, PM history, failures, telemetry, and business impact
Analytics Layer	PM-to-failure linker producing days_to_next_failure per PM event
ML Layer	Logistic Regression + Random Forest classifiers for 14d and 30d failure horizons
Prescriptive Layer	Risk-band action planner combining ML scores, PM history, and financial exposure
UI Layer	Streamlit 9 tabs: executive summary → telemetry → ML training → data export
Deployment Layer	Docker image for HF Spaces · Snowflake Streamlit · Local pip install

Dependency Stack

```
streamlit>=1.28,<2.0
pandas>=2.0
numpy>=1.24
scikit-learn>=1.3
plotly>=5.18
matplotlib>=3.7
openpyxl>=3.1
python-dateutil>=2.8
```

Containerization

```
FROM python:3.11-slim

RUN apt-get update && apt-get install -y --no-install-recommends \
    build-essential \
    && rm -rf /var/lib/apt/lists/*
```

```

RUN useradd -m -u 1000 user
USER user
ENV PATH="/home/user/.local/bin:$PATH"
WORKDIR /home/user/app

COPY --chown=user requirements.txt requirements.txt
RUN pip install --no-cache-dir --upgrade pip && \
    pip install --no-cache-dir -r requirements.txt

COPY --chown=user . .

EXPOSE 7860

CMD ["streamlit", "run", "streamlit_app.py", \
    "--server.port=7860", \
    "--server.address=0.0.0.0", \
    "--server.headless=true", \
    "--server.enableCORS=false", \
    "--server.enableXsrfProtection=false", \
    "--browser.gatherUsageStats=false"]

```

Deployment Path

```

flowchart LR
    A[Local Dev · pip install + streamlit run] --> B[Synthetic dataset auto-loaded]
    B --> C[Docker build]
    C --> D[Push to Hugging Face Spaces]
    D --> E[Public HTTPS endpoint]
    E --> F[CSV upload or demo data mode]
    E --> G[Live ML training in-browser]
    E --> H[Downloadable risk reports]
    B --> I[Snowflake Streamlit]
    I --> J[environment.yml runtime]
    J --> K[Warehouse-backed data queries]

```

Runtime Example

```

# Clone the repo
git clone
https://github.com/dranubhaparashar/Predictive_Preventive_Maintenance_for_Generator_Reliability.
git
cd Predictive_Preventive_Maintenance_for_Generator_Reliability

```

```
# Create virtual environment
python -m venv pm_env
source pm_env/bin/activate

# Install dependencies
pip install -r requirements.txt

# Run the app
streamlit run streamlit_app.py --server.port 8505

# Open http://localhost:8505
# Auto-loads synthetic generator dataset
# Navigate to "Predictive ML & PM Strategy" tab to train models
# Export risk rankings from the Data Explorer tab
```

Demo Video

Embedded media: [Open video](#)

Public Demo Surfaces

Important The Hugging Face Space provides instant browser-based access to the full dashboard with the embedded synthetic dataset — no local install required. The GitHub repository contains source code, architecture docs, and the function reference wiki.

Demo Links

- 🚀 **Live App:** [AnubhaParashar / Predictive Preventive Maintenance for Generator Reliability](#)
 - 📁 **GitHub Repository:** [Predictive Preventive Maintenance for Generator Reliability](#)
 - 📖 **Wiki Documentation:** [Architecture](#) · [Function Reference](#) · [Technical Spec](#)
-

Engineering Value

This project is not just about predicting failures.

It demonstrates how to take raw maintenance records and sensor data and make them:

- **measurable** — PM effectiveness expressed as a single comparable metric per event
- **predictive** — binary classifiers producing per-asset failure probabilities with calibrated horizons
- **prescriptive** — risk scores translated into prioritized action recommendations with business-impact weighting

- **interactive** — 9-tab Streamlit UI that works with any CSV upload or the embedded demo data
 - **containerized** — Docker image deployable to Hugging Face, Snowflake, or any container host
 - **publicly demoable** — HF Space with synthetic dataset and live model training
-

Current Strengths

- PM-to-failure linker produces a deterministic, interpretable reliability metric without ML
 - On-time vs. delayed PM comparison quantifies the operational cost of deferred maintenance
 - Two-model approach gives both interpretability (LR coefficients) and accuracy (RF feature importance)
 - Dual prediction horizons (14d and 30d) support both urgent dispatch and weekly planning cycles
 - Prescriptive action layer closes the loop from prediction to work-order recommendation
 - Multi-platform deployment (HF Spaces, Snowflake, local) accommodates different production constraints
 - Synthetic dataset is structurally identical to real production schemas — swap files to go live
-

Next Improvements

- Connect real PM work-order systems (Maximo, ServiceNow, SAP PM) via API
 - Add survival analysis (Kaplan-Meier, Cox regression) alongside classification models
 - Implement scheduled model retraining with labeled failure outcome ingestion
 - Add alerting integration: email or Slack notification on Critical risk assets
 - Add multi-asset comparison dashboard for fleet-level reliability benchmarking
 - Add time-series anomaly detection (LSTM or Isolation Forest) on raw telemetry streams
 - Add Snowflake-native scheduled refresh with Streamlit Tasks
 - Add confidence intervals on ML predictions for uncertainty-aware decision making
 - Add A/B testing framework for PM interval variants
 - Add geospatial map view for regional failure density analysis
-

Key Innovation

Important This project connects **PM effectiveness measurement** with **predictive ML risk scoring** and **prescriptive action planning** — three layers that maintenance teams typically operate separately in disconnected spreadsheets and dashboards.

It turns:

Raw maintenance logs → Survival duration per PM event → Failure probability per asset → Prioritized maintenance action per generator

rather than stopping at failure reporting.

The `days_to_next_failure` linker is the keystone — it shows that **operational history (what happened after each PM) is just as valuable as sensor telemetry** when building a reliability model.

Conclusion

This repository shows how generator maintenance analytics can evolve from reactive failure tracking into a predictive reliability workflow with prescriptive recommendations.

It combines:

- Deterministic PM-to-failure event linking with `days_to_next_failure` as the core metric
 - Supervised ML classifiers (Logistic Regression + Random Forest) for 14/30-day failure horizons
 - Prescriptive action planner merging risk scores with business-impact data
 - 9-tab Streamlit UI covering the full analytical journey from raw data to maintenance actions
 - Synthetic dataset mirroring real production schemas for immediate portability
 - Docker deployment to Hugging Face Spaces and Snowflake Streamlit
 - Comprehensive wiki documentation covering architecture and function references
-

Final Thought

From **counting failures** to **predicting them before they happen — and telling you what to do about it**

The real value is not only the Random Forest classifier — it is the **PM effectiveness metric, the prescriptive action layer, and the business-impact weighting** that together make the system operationally useful rather than just analytically interesting.

9 dashboard tabs · 2 ML models · 2 failure horizons · 7 input datasets · 1 clear recommendation per generator.

